


# ***The Illusion Of Execution***




Nitsan Wakart (@nitsanw)  
Lead Performance Engineer, Azul Systems

Thanks!

# I work on Zing!

- Awesome JVM, will be mentioned in slides
- Only on Linux/x86
- Aimed at server side systems
- Highly focused on responsiveness



# PSYCHOSOMATIC, LOBOTOMY, SAW

It's X, you'll need Y, I'll get Z

Monday, 1 December 2014

## *The Escape of ArrayList.iterator()*

{This post assumes some familiarity with JMH. For more JMH related content start at the new and improved **JMH Resources Page** and branch out from there!}

**Escape Analysis** was a much celebrated optimisation added to the JVM in Java 6u23:

"Based on escape analysis, an object's escape state might be one of the following:

- **GlobalEscape** – An object escapes the method and thread. For example, an object stored in a static field, or, stored in a field of an escaped object, or, returned as the result of the current method.
- **ArgEscape** – An object passed as an argument or referenced



<http://psy-lob-saw.blogspot.com>



FORK ME ON G

# JCTOOLS

Java Concurrency Tools

download  
**.ZIP**

download  
**.TGZ**

## *JCTools - Java Concurrency Tools for the JVM*

This project aims to offer some concurrent data structures currently missing from the JDK:

- › SPSC/MPSC/SPMC/MPMC Bounded lock free queues
- › SPSC/MPSC Unbounded lock free queues
- › Alternative interfaces for queues (experimental)
- › Offheap concurrent ring buffer for JTC/IPC purposes (experimental)


<https://github.com/JCTools/JCTools>

# The JVM is a Magical Place

- You write some Java/Scala/Clojure/JavaScript/Ruby
- Compile/Pack/Deploy

BOOM!

- JVM “executes” the “code”
- Use infinite memory!
- Run infinite threads!



# The Way Down

Your (byte) code...

Is executed (interpreter/C1/C2)...


by a JVM (OpenJDK/Oracle/IBM/Zing)...

Which is a process of an OS (specific native libraries)...

Which is running on some hardware (specific instruction set)...

```
public void foo(Bar bar) {  
    int nogCount = 0;  
    for (int i = 0; i < 10; i++) {  
        if (bar.getZog(i).isNog()) nogCount++;  
    }  
    while (nightIsYoung) {  
        nogCount += hit(bar);  
    }  
    if (nogCount > MAX_NOG)  
        throw new NogOverflowError();  
}
```





# Why Stop The World?

- Some GC phases
- Deoptimization
- Stack trace dump (and other JVMTI activities)
- Lock un-biasing
- Class redefinition
- And more!

See excellent talks:

<https://www.youtube.com/watch?v=Y39klIzX1P8> : “With GC Solved, What Else Makes a JVM Pause?” by John Cutherson

<https://vimeo.com/120533011> : “When Does the JVM JIT & Deoptimize?” by Doug Hawkins

Can one Java thread  
stop all other Java  
threads?


# You too can trigger a STW pause!

- On normal allocation (Young Gen full)
- On large object allocation (Old Gen full)
- On synchronized block (unbiasing)
- Hitting cold code/loading new classes (deoptimization)
- Profiling (JVMTI *GetStackTrace()*)



# Zing: The pause-less JVM

- C4 - Fully Concurrent GC (including young gen)
- ReadyNow! - Persistent code profile, fighting Deopts
- No biased locking
- JVMTI *GetStackTrace* does not STW
- ... some STW pauses exist (smaller than OS hiccups)



# The JVM....





... Must give  
us pause.

# jHiccup: A STW Pause Detector

- Try to run every X milliseconds
- Record deviations from schedule
- Make lovely graphs!

See: <https://github.com/giltene/jHiccup>



# How does a JVM stop the world?

To 'Stop The World' the JVM brings all **Mutator** threads to a **Safepoint**

- Raise **Safepoint** 'flag'
- Wait for ALL threads to reach **Safepoint** state



Safe == Good

Safepoints == Goodpoints?

# At a Safepoint

“...the thread's representation of it's Java machine state is well described, and can be safely manipulated and observed by other threads in the JVM”

See full quote from Gil Tene, originally posted on the “Mechanical Sympathy” mailing list:  
<http://chriskirk.blogspot.ru/2013/09/what-is-java-safepoint.html>

# Safepoint

(noun.)

A JVM thread state

- Waiting/Idle/Blocked → @Safepoint
- Running Java code → !@Safepoint
- Running JNI code → @Safepoint

<http://blog.ragozin.info/2012/10/safepoints-in-hotspot-jvm.html>

<http://psy-lob-saw.blogspot.com/2014/03/where-is-my-safepoint.html>



# Where do we see a Safepoint poll?

- Between every 2 bytecodes (interpreter)
- Backedge of non-'counted' loops (C1/C2)
- Method exit (C1/C2)
- JNI call exit

```
public void foo(Bar bar) {  
    int nogCount = 0;  
    for (int i = 0; i < 10; i++) {  
        if (bar.getZog(i).isNog()) nogCount++;  
    }  
    while (nightIsYoung) {  
        nogCount += hit(bar);  
    }  
    if (nogCount > MAX_NOG)  
        throw new NogOverflowError();  
}
```

```
public void foo(Bar bar) {  
    int nogCount = 0;  
    for (int i = 0; i < 10; i++) {  
        if (bar.getZog(i).isNog()) nogCount++;  
    }  
    while (nightIsYoung) {  
        nogCount += hit(bar);  
        // Safepoint poll  
    }  
    if (nogCount > MAX_NOG)  
        throw new NogOverflowError();  
    // Safepoint poll  
}
```

# Safepoint poll: OpenJDK

- Read from a special page:

```
"test    DWORD PTR [rip+0xffffffffffe690e53],eax"
```

- JVM Sets the page to protected, polling threads trap a SEGV and go to safepoint
- Look for `{poll}` or `{poll_return}` in the assembly comments



# Safepoint poll: Zing

- Read the thread local safe point flag:

```
"gs:cmp4i [0x40 tls._please_self_suspend],0  
jnz 0x500a0186; Where the safepoint code be"
```

- JVM Sets the thread flag to 1, polling threads hop to
- Look for **"tls.\_please\_self\_suspend"**

```
public void foo(Bar bar) {  
    int nogCount = 0;  
    for (int i = 0; i < 10; i++) {  
        if (bar.getZog(i).isNog()) nogCount++;  
    }  
    while (nightIsYoung) {  
        nogCount += hit(bar);  
        if (MUST_SAFEPOINT) gotoSafepoint();  
    }  
    if (nogCount > MAX_NOG)  
        throw new NogOverflowError();  
    if (MUST_SAFEPOINT) gotoSafepoint();  
}
```

Safepoint polls == Overhead



It's just a harmless lil' safepoint they said

# How Bad?

```
private static final class DataSet {  
    private final int[] data;  
  
    int intSize() { return data.length; }  
    int intGet(int index) { return data[index]; }  
    void intSet(int index, int v) { data[index] = v; }  
  
    long longSize() { return data.length; }  
    int longGet(long index) { return data[(int) index]; }  
    void longSet(long index, int v) { data[(int) index] = v; }  
}
```

See: <https://github.com/netty/netty/pull/3969#issuecomment-132559757>

```
@Benchmark
public int sumNoSafePoint() {
    int sum = 0;
    for (int index = 0; index < datasetA.intSize(); ++index) {
        sum += datasetA.intGet(index);
    }
    return sum;
}
```

```
@Benchmark
public int sumSafePoint() {
    int sum = 0;
    for (long index = 0; index < datasetA.longSize(); ++index) {
        sum += datasetA.longGet(index);
    }
    return sum;
}
```

# How Bad?

| Benchmark                   | Score     | Error   | Units |               |
|-----------------------------|-----------|---------|-------|---------------|
| Benchmark.copyNoSafepoint   | 76.943 ±  | 14.256  | ns/op |               |
| Benchmark.copySafepoint     | 796.583 ± | 55.583  | ns/op | ← 10x worse!  |
| Benchmark.equalsNoSafepoint | 367.192 ± | 16.398  | ns/op |               |
| Benchmark.equalsSafepoint   | 806.421 ± | 196.106 | ns/op | ← 2.1x worse! |
| Benchmark.fillNoSafepoint   | 84.075 ±  | 19.033  | ns/op |               |
| Benchmark.fillSafepoint     | 567.866 ± | 10.154  | ns/op | ← 7x worse!   |
| Benchmark.sumNoSafepoint    | 338.204 ± | 44.529  | ns/op |               |
| Benchmark.sumSafepoint      | 585.657 ± | 105.808 | ns/op | ← 1.5x worse! |

# Safepoint poll side effects

- Can prevent loop unrolling
- Can prevent SuperWord optimizations
- Can prevent OptimizeFill optimization
- ... Inhibits the compiler



Safepoint polls == BAD!

DEMO: When will it exit?

Delaying a Safepoint == BAD!

# How long was this GC pause?

```
[GC [PSYoungGen: 109884K->14201K(139904K)]  
691015K→595332K(1119040K),  
0.0454530 secs]
```

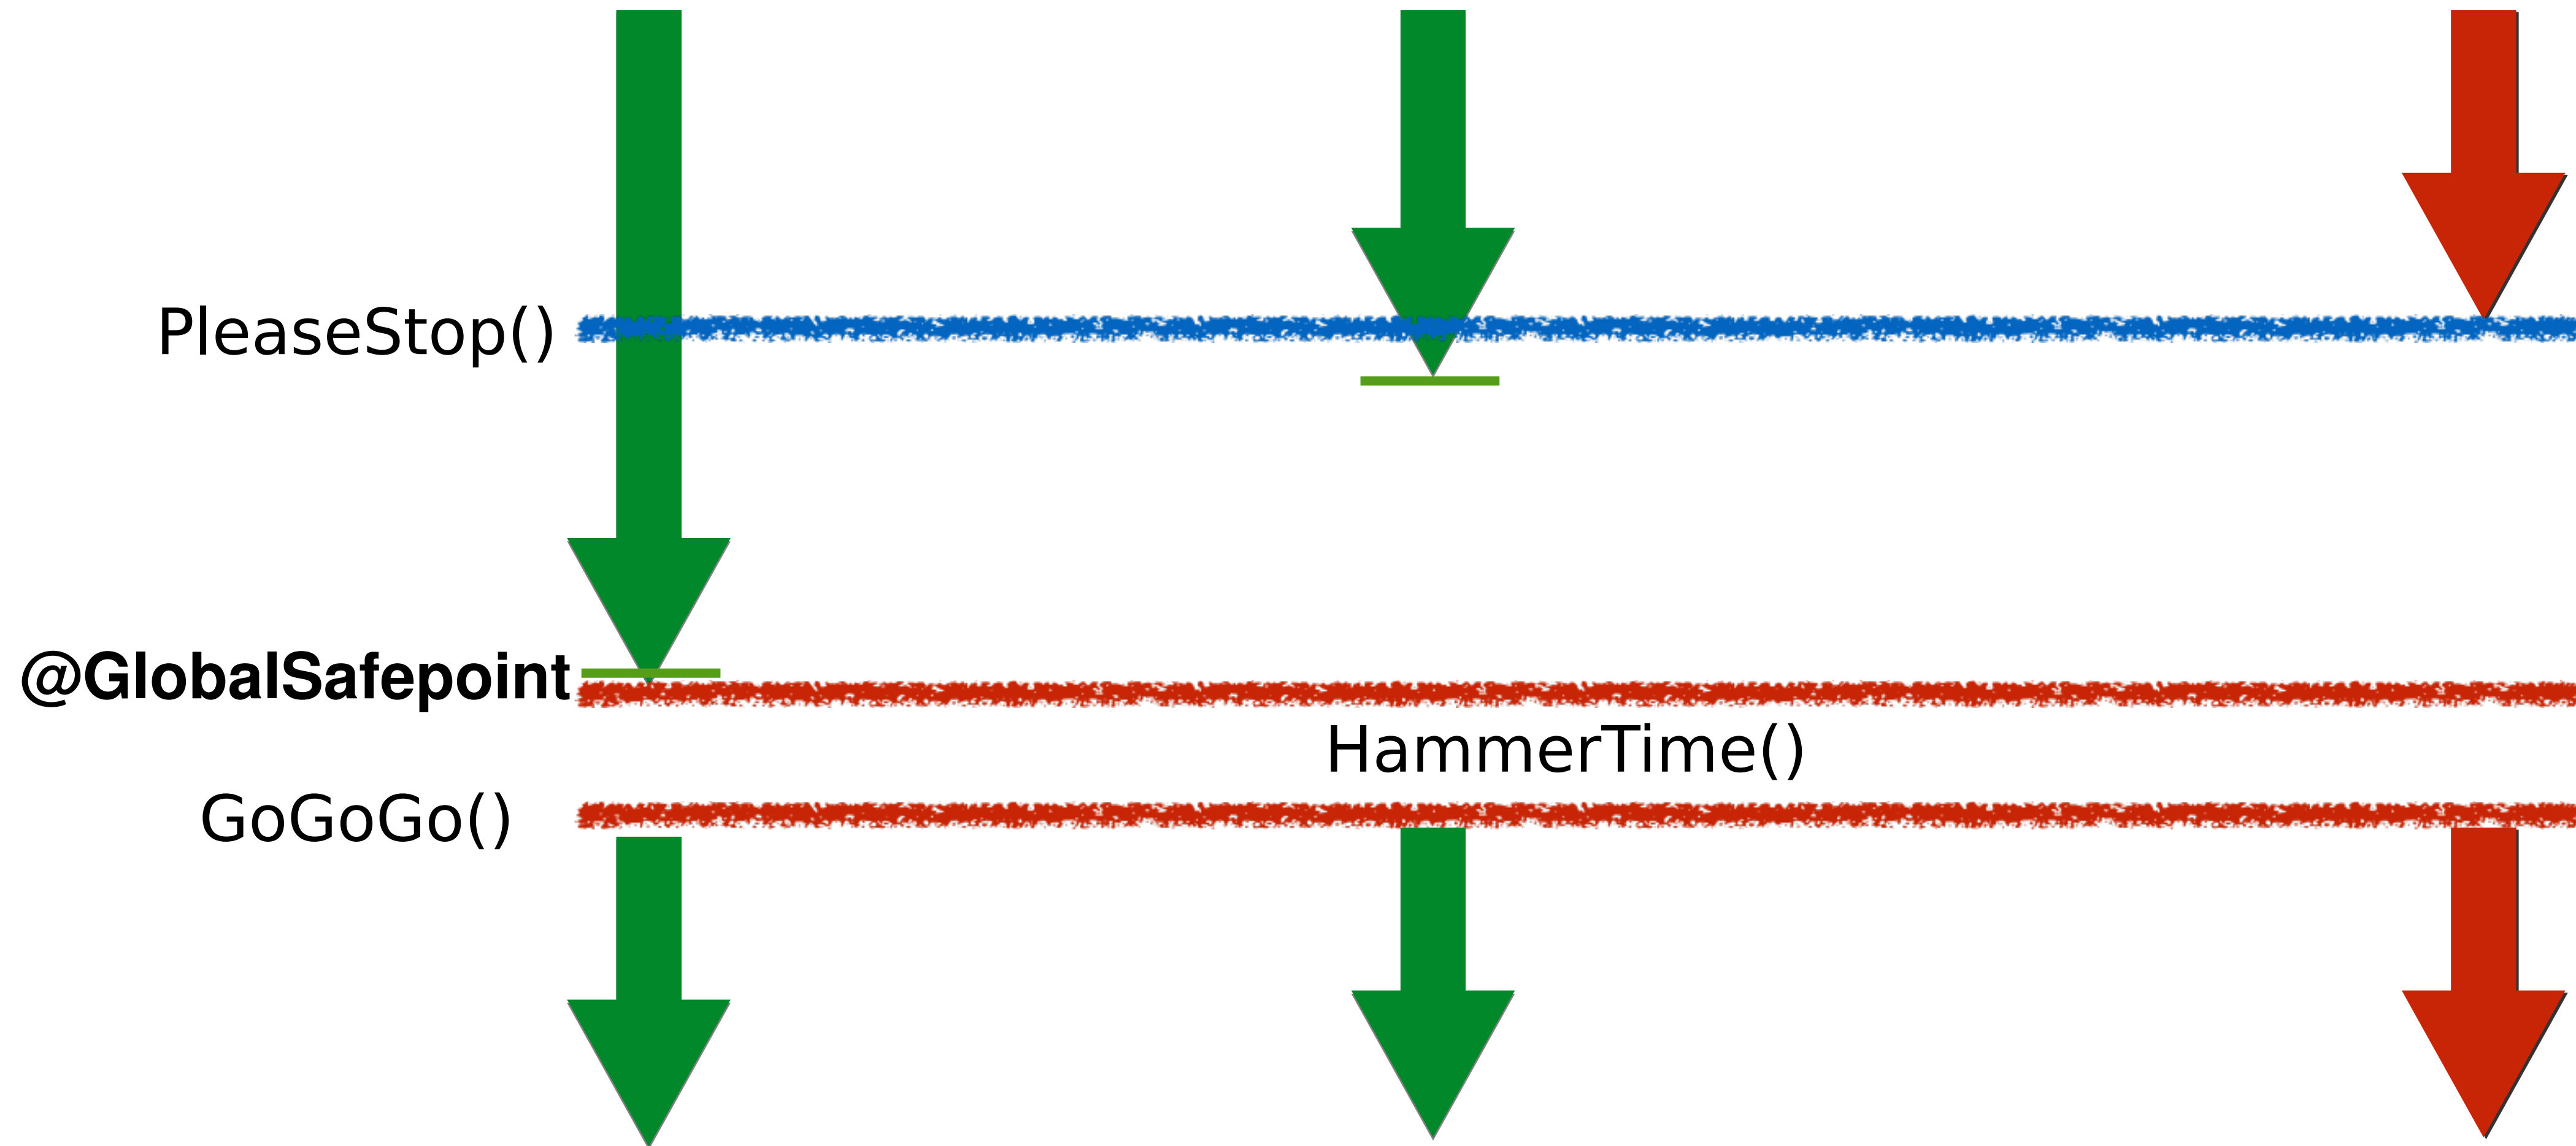
# Thread States Legend

**Waiting - Blocked (IO/wait/park)**

**Ready - Queued for execution**

**Running – On CPU**

# TTSP - Time To Safepoint



# Where is TTSP reported?

- Zing: Reported as part of pause time (with break down)
- OpenJDK: Not included in GC pause time
- OpenJDK: Must use `-XX:+PrintGCApplicationStoppedTime`

Total time for which application threads were stopped: **4.2399590** seconds, Stopping threads took: **4.2398886** seconds


- OpenJDK: `-XX:+PrintSafepointStatistics -XX:PrintSafepointStatisticsCount=1`

**Max Actual Pause = TTSP + work in safepoint**

Long TTSP == BAD!



Safepoint? hmm yes...  
Poll it you must



# So what?

- TTSP can dominate pause times, track it!
- No safepoint polls in counted loops:
  - ✓ On Zing can use `-XX+KeepSafepointsInCountedLoops`
- Some 'big' operations may have no safepoint polls:
  - x `System.arraycopy` (chunked on Zing)
  - x `Object clone/init` (chunked on Zing)
  - x `String.indexOf`
  - x Others...



Time To Safepoint  
vs.  
Safepoint poll overhead

# Safepoint poll insertion

- Backedge of non-'counted' loops
- Method exit (removed by inlining)

**Heuristics based on instruction count**

Instruction count  $\neq$  Time

Is it faster to read  
from a HashMap or  
a file?



I Like to Move it...Move it..

He Likes to Move it...Move it..

You Like To.....?

**MOVE IT !**

# Why is MOV important?

mov r10d, DWORD PTR [rsi+0x10] ; \*getfield hold/Holder::set@2 (line 8)

mov r11, r10

shl r11, 0x3

xor eax, eax

cmp rdx, r11

je 0x000000010b065eb2 ; Holder::set@5 (line 8)

mov r8d, DWORD PTR [rdx+0x8] ; implicit exception: dispatches to 0x000000010b065f11

cmp r8d, 0xef5d495b ; {oop('java/lang/Integer')}

jne 0x000000010b065ebe ; \*invokevirtual equals/Holder::set@ (line 11)

mov r8d, DWORD PTR [r12+r10\*8+0x8] ; implicit exception: dispatches to 0x000000010b065ef9

cmp r8d, 0xef5d495b ; {oop('java/lang/Integer')}

jne 0x000000010b065edd ; \*instanceof/Integer::equals@ (line 764)/Holder::set@ (line 11)

mov r8d, DWORD PTR [rdx+0xc]

shl r10, 0x3 ; \*checkcast/Integer::equals@ (line 765)/Holder::set@ (line 11)

mov r11d, DWORD PTR [r10+0xc]

cmp r8d, r11d

je 0x000000010b065eb2 ; - java.lang.Integer::equals@18 (line 765)/Holder::set@ (line 11)

inc DWORD PTR [rsi+0xc] ; \*putfield setCount/Holder::set@ (line 15)

mov r10, rsi

mov r11, rdx

shr r11, 0x3

mov DWORD PTR [rsi+0x10], r11d

shr r10, 0x9

movabs r11, 0x10a4e9000

mov BYTE PTR [r11+r10\*1], r12b ; \*putfield hold/Holder::set@ (line 14)

mov eax, 0x1

add rsp, 0x30; This is 0x000000010b065eb2

pop rbp

test DWORD PTR [rip+0xfffffffffff71c143], eax # 0x000000010a782000; {poll\_return}

ret



# Memory Topology



Converting to cycles, assumed 3 cycles per nano-second

[http://www.eecs.berkeley.edu/~rcs/research/interactive\\_latency.html](http://www.eecs.berkeley.edu/~rcs/research/interactive_latency.html)


# Beer Cache Hierarchy

- Register - Bottle at your lips (1 sec)
- L1/L2 - ... on table (4 - 10 secs)
- LLC - ... waiting on the bar (40 secs)
- Main memory - ... at store across the road (5 mins)
- SSD - ... in another country? (13.3 hrs)
- HD - ... on the moon? (34 days, 17 hrs)

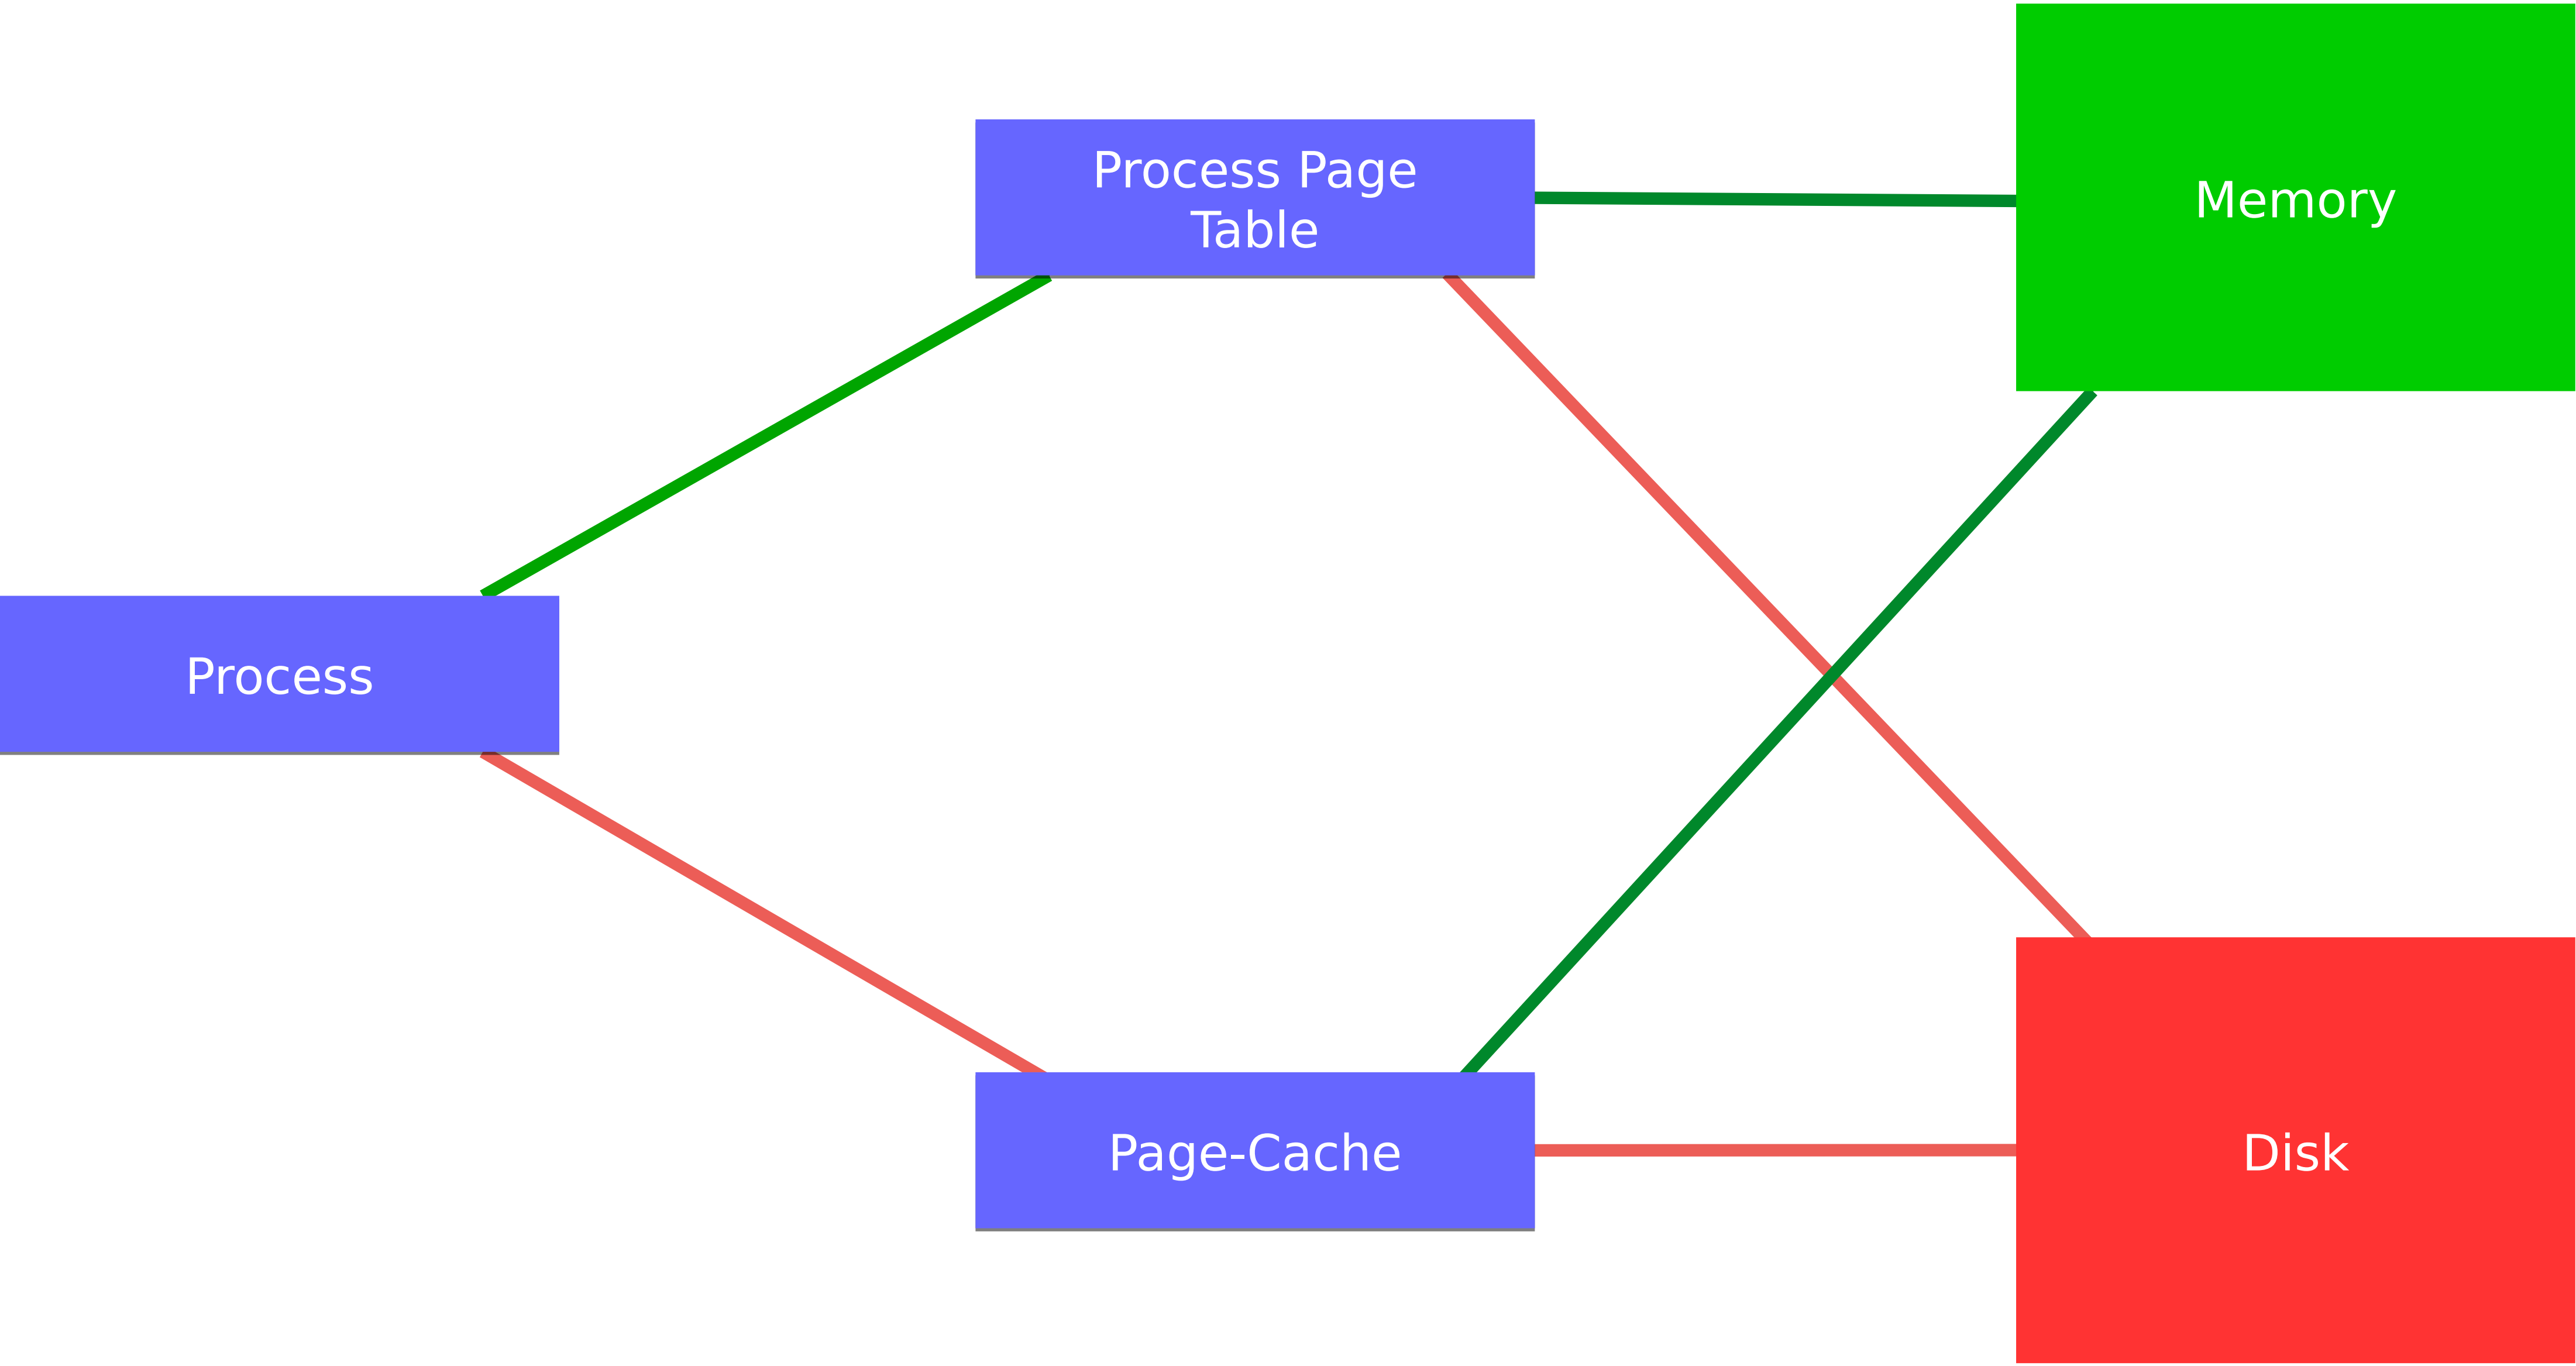
See: <http://blog.netopyr.com/2014/11/28/reactions-to-the-beer-cache-hierarchy/>

# Resident vs. Virtual Memory

- How much memory can a process use?
- OS maps when actually used
- What happens when we exceed physical?
- Swap
- Page faults



"Resident evil ver4" by impawards. Licensed under Fair use of copyrighted material in the context of Resident Evil (film) via Wikipedia - [http://en.wikipedia.org/wiki/File:Resident\\_evil\\_ver4.jpg#mediaviewer/File:Resident\\_evil\\_ver4.jpg](http://en.wikipedia.org/wiki/File:Resident_evil_ver4.jpg#mediaviewer/File:Resident_evil_ver4.jpg)



# So What?

- Memory access assumed short (1 instruction)
- Page fault → very large TTSP
- `MappedByteBuffer` → IO out of safepoint, mostly great...  
but see prev. point
  - ✓ Use `mlock` (requires JNI call, not part of JDK)?



# So What?

- Benchmarking? Use relevantly sized data sets and access patterns
- Avoid using more memory than available
  - ✓ disable swap? (Set swappiness=0)
  - ✓ Buy more memory
- Monitor Page faults
- Consider priming memory?



Have you ever written a  
single threaded Java  
application?



# The JVM Process: Threads

How many threads for an application?

- Application Java threads (Main, Thread etc)
- Application native threads (native lib)
- JVM Threads (GC, Compiler, JMX, RMI...)

**There's no such thing as a single threaded Java application**

# Threads Example

(Oracle JDK8, on i5/dual core laptop, no args)

- Application threads
- 4 GC Threads (ParallelGC)
- 3 Compiler threads (1 C1 + 2 C2)
- ... and others
- 15 threads reported by jstack, 19 by OS

# Multi-Tasking OS

More processes than cores

- Yay!!! Gimme ALL the threads!
- ... unless they all want to run at the same time

Scheduling and interrupts

- Fairness (thread priority, starvation)
- Context Switching (overhead)

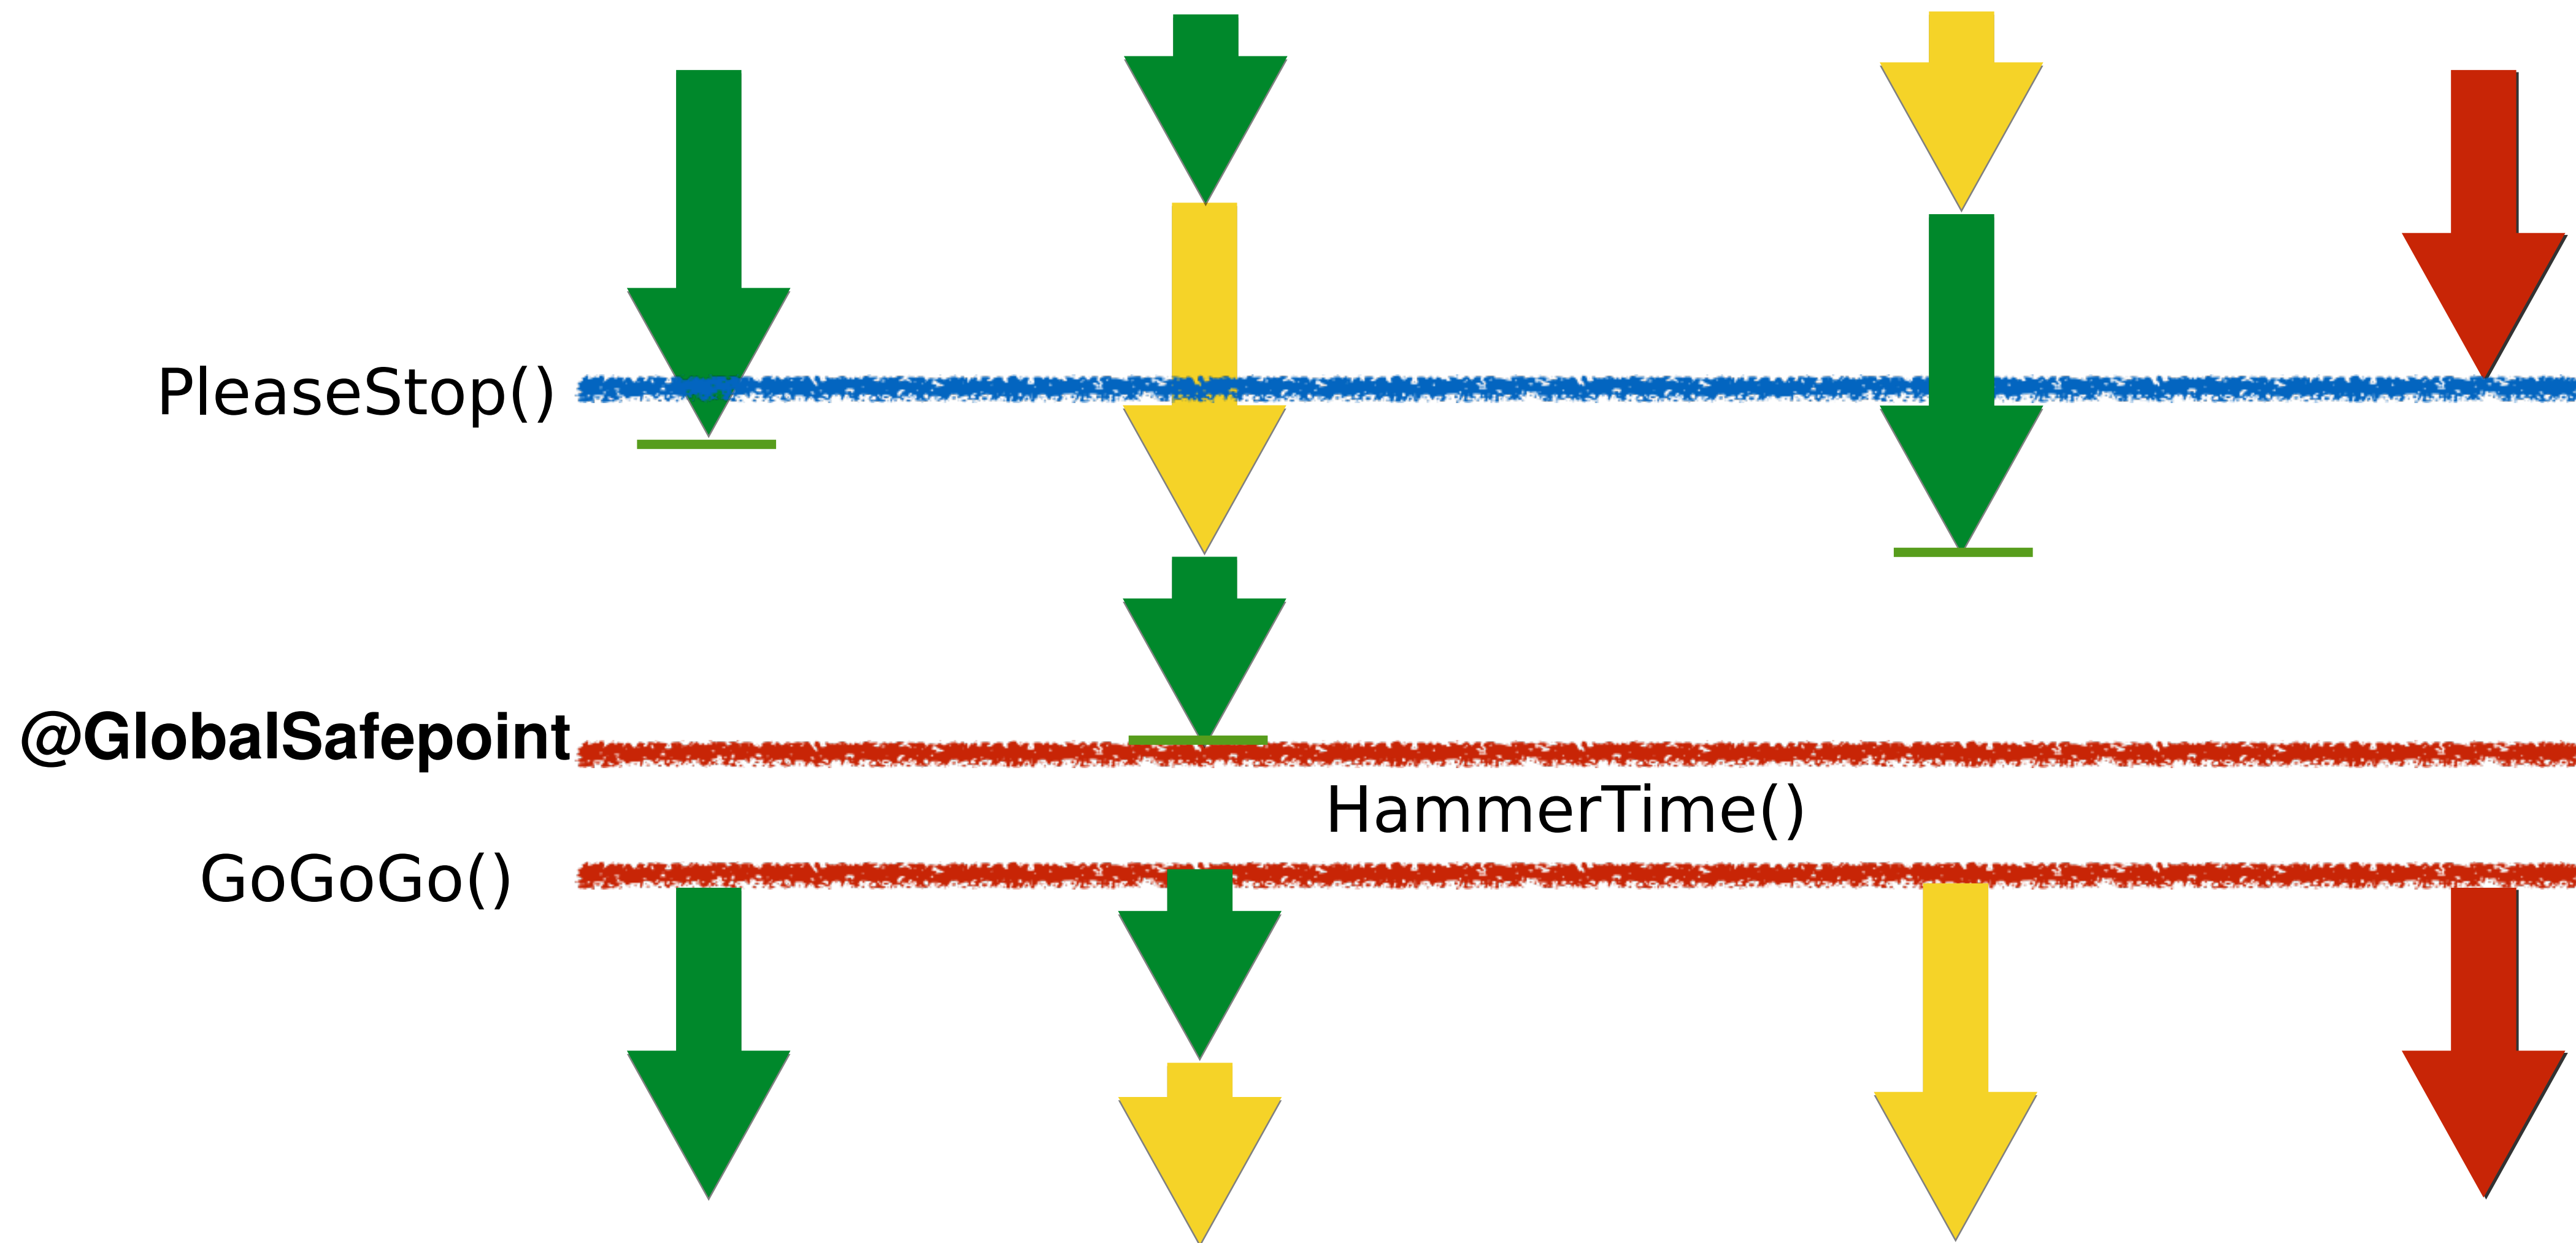


# So What?

- OS can suspend Java threads
- Threads suspension – not necessarily @Safepoint
- All suspended threads must resume and reach safepoint before safepoint work can start
- Runnable Threads > Cores → Longer TTSP



# TTSP - Time To Safepoint



# So What?

- Try and avoid having more **ready/running** threads than cores
  - ✓ Monitor run queue (avg. load), ctx switches
- Consider controlling OS resource management
  - ✓ Use taskset/numactl
  - ✓ JVM is not aware of taskset/numactl/isocpus
- Consider controlling JVM thread counts




# OOPs

hold = love;




```
mov    r10,rsi
mov    r11,rdx
shr    r11,0x3
mov    DWORD PTR [rsi+0x10],r11d
shr    r10,0x9
movabs r11,0x10a4e9000
mov    BYTE PTR [r11+r10*1],r12b
```

Compressed Oops



Card Marking



# What's an OOP?

- Ordinary Object Pointer
- Java: Object reference -> JVM: OOP
- Pointers to managed data on the heap



# Java Hidden Symbols

```
void copyPoint(Point p1, Point p2) {  
    p1.x = p2.x;  
}
```

```
void copyPoint(oop p1, oop p2) {  
    address a1 = readBarrier(p1);  
    address a2 = readBarrier(p2);  
    oop x = getObject(a2 + xFieldOffset);  
    putObject(a1 + xFieldOffset, x);  
    writeBarrier(a1, x);  
    safepoint_poll();  
}
```

# Memory Barrier

(not the JMM kind)

“...a block [of code, executed] on reading from or writing to certain memory locations by certain threads or processes.”

Memory Management Reference:

<http://www.memorymanagement.org/glossary/b.html#term-barrier-1>

# Compressed OOPs

-XX:+UseCompressedOops

- Want large heaps ( $> 4\text{Gb}$ ) and 32bit addresses
- Objects aligned on 8 bytes\*
- Can compress OOP by dropping last 3 bits ( $>>3$ )
- Must decompress address to use it ( $<<3$ )
- Max referenced heap size is now  $4\text{Gb} * 8 = 32\text{Gb}$

<https://wikis.oracle.com/display/HotSpotInternals/CompressedOops>

\* Can change with -XX:ObjectAlignmentInBytes

# Compressed Oops

## Example(x86):

JAVA:

```
long v = this.l.longValue();
```

-XX:-UseCompressedOops:

```
mov r10, QWORD PTR [rsi+0x18]      ; r10= this.l  
mov r10, QWORD PTR [r10+0x10]     ; r10= l.value
```

-XX:+UseCompressedOops:

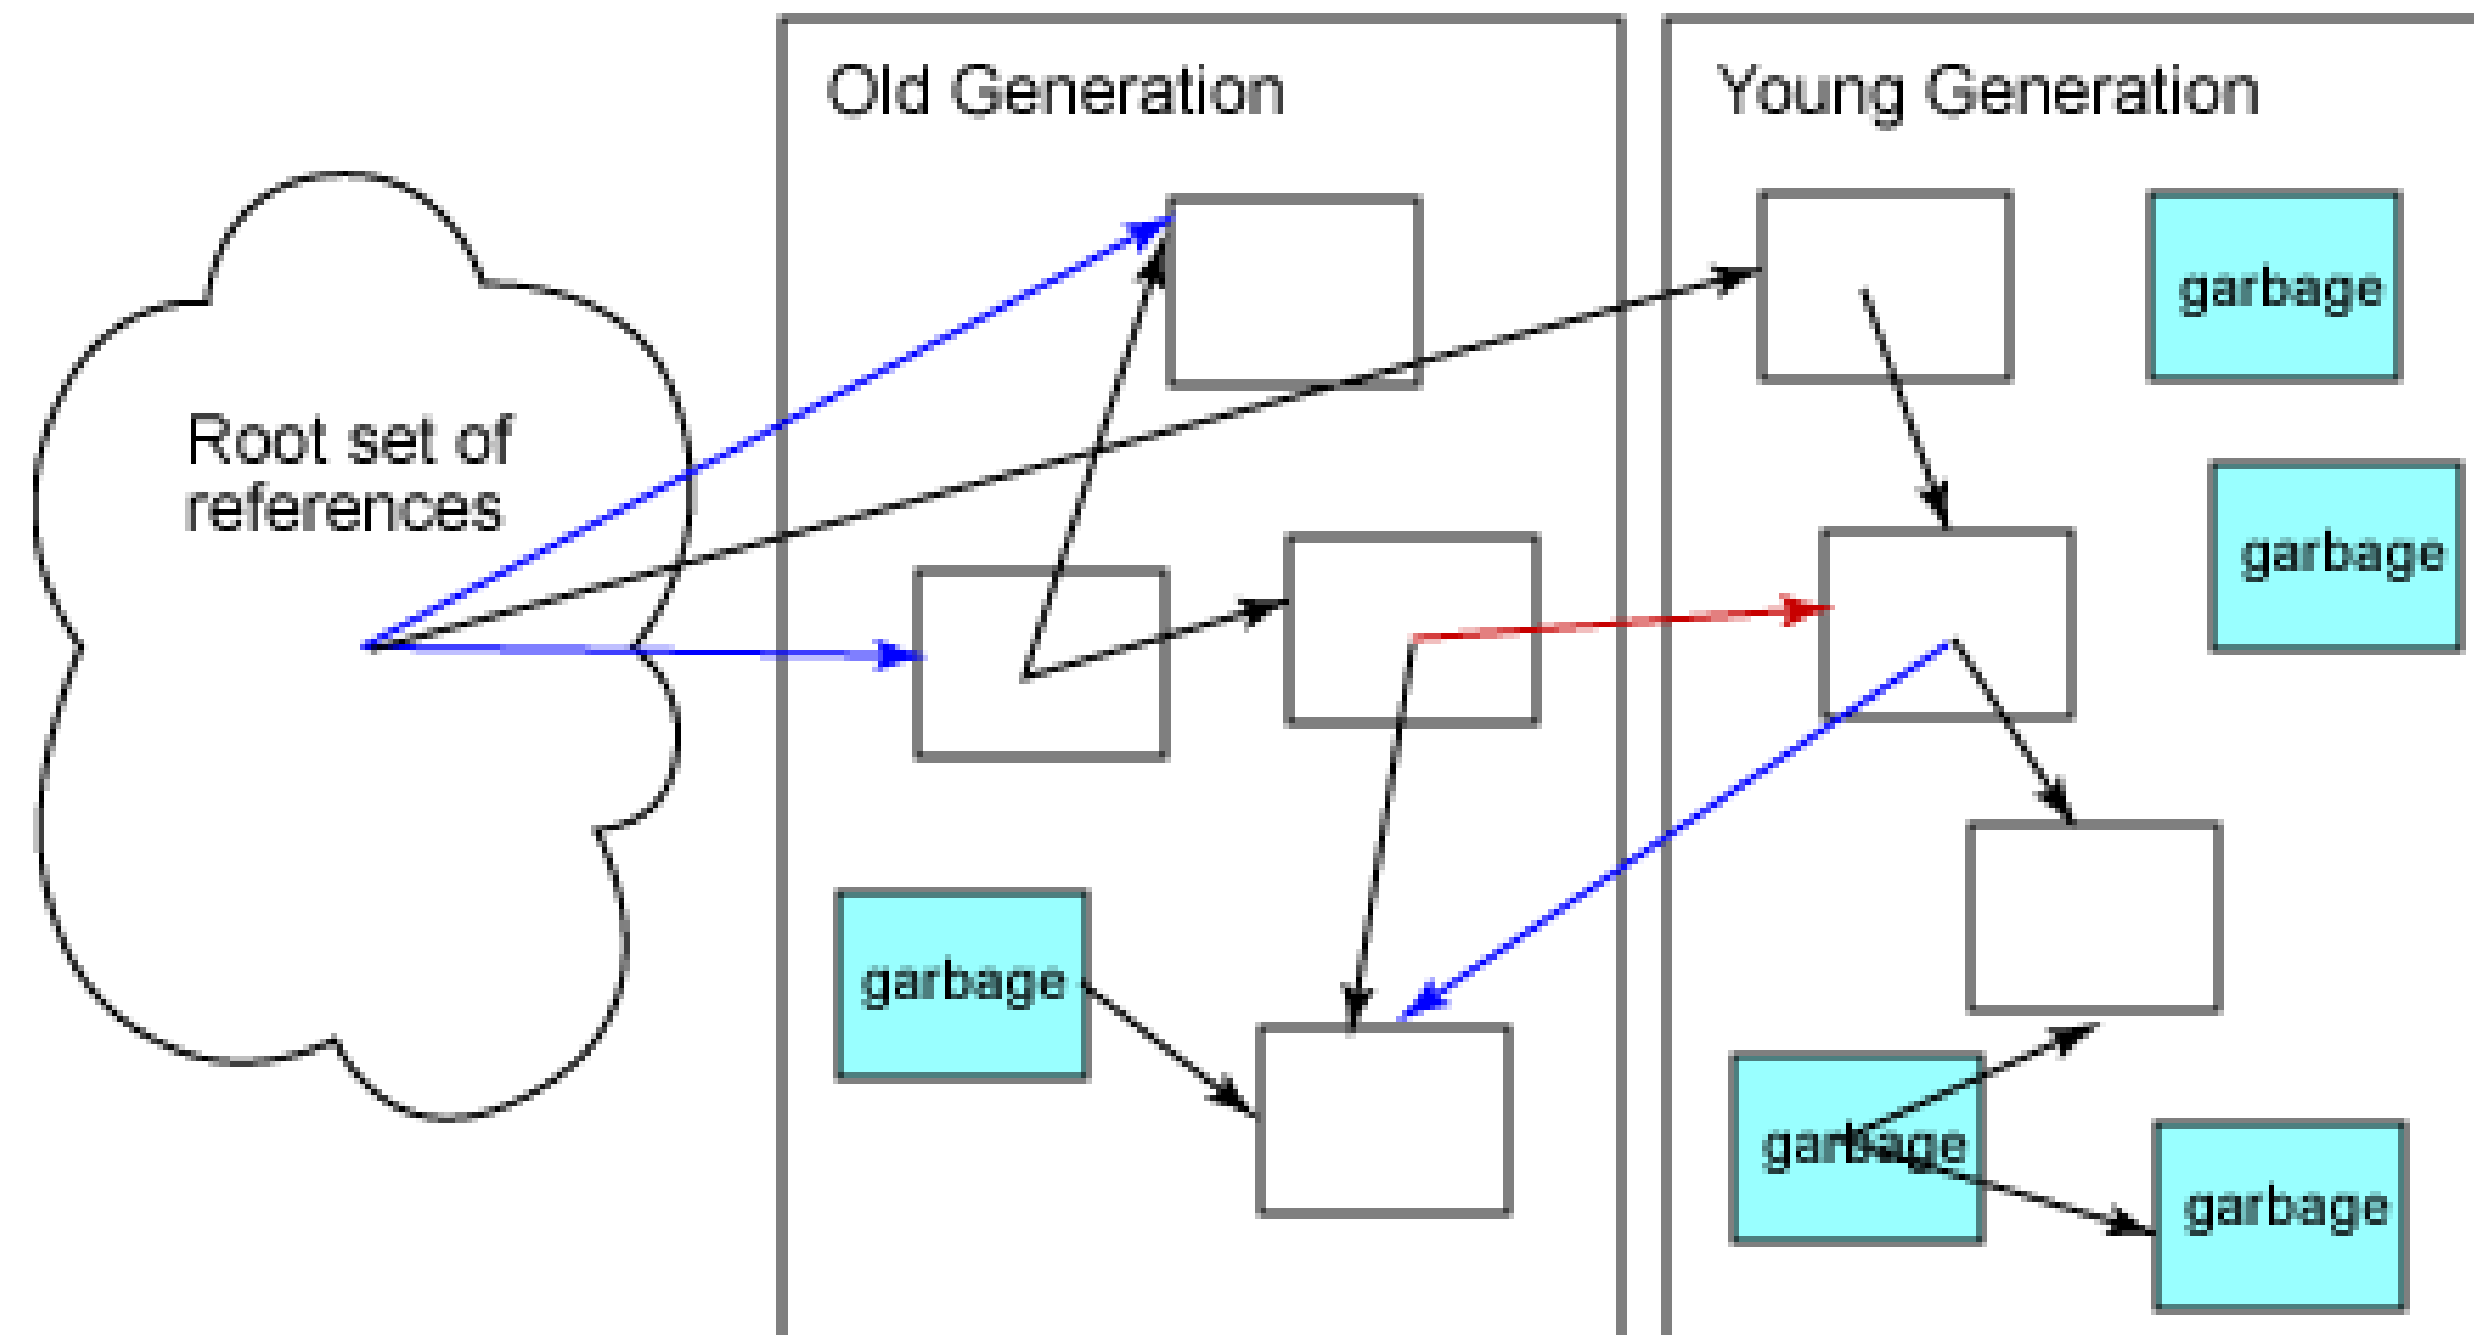
```
mov r11d, DWORD PTR [rsi+0x10]     ; r11d= this.l  
mov r10, QWORD PTR [r12+r11*8+0x10] ; r10= l.value
```

# CompressedOops instanceof ReadBarrier

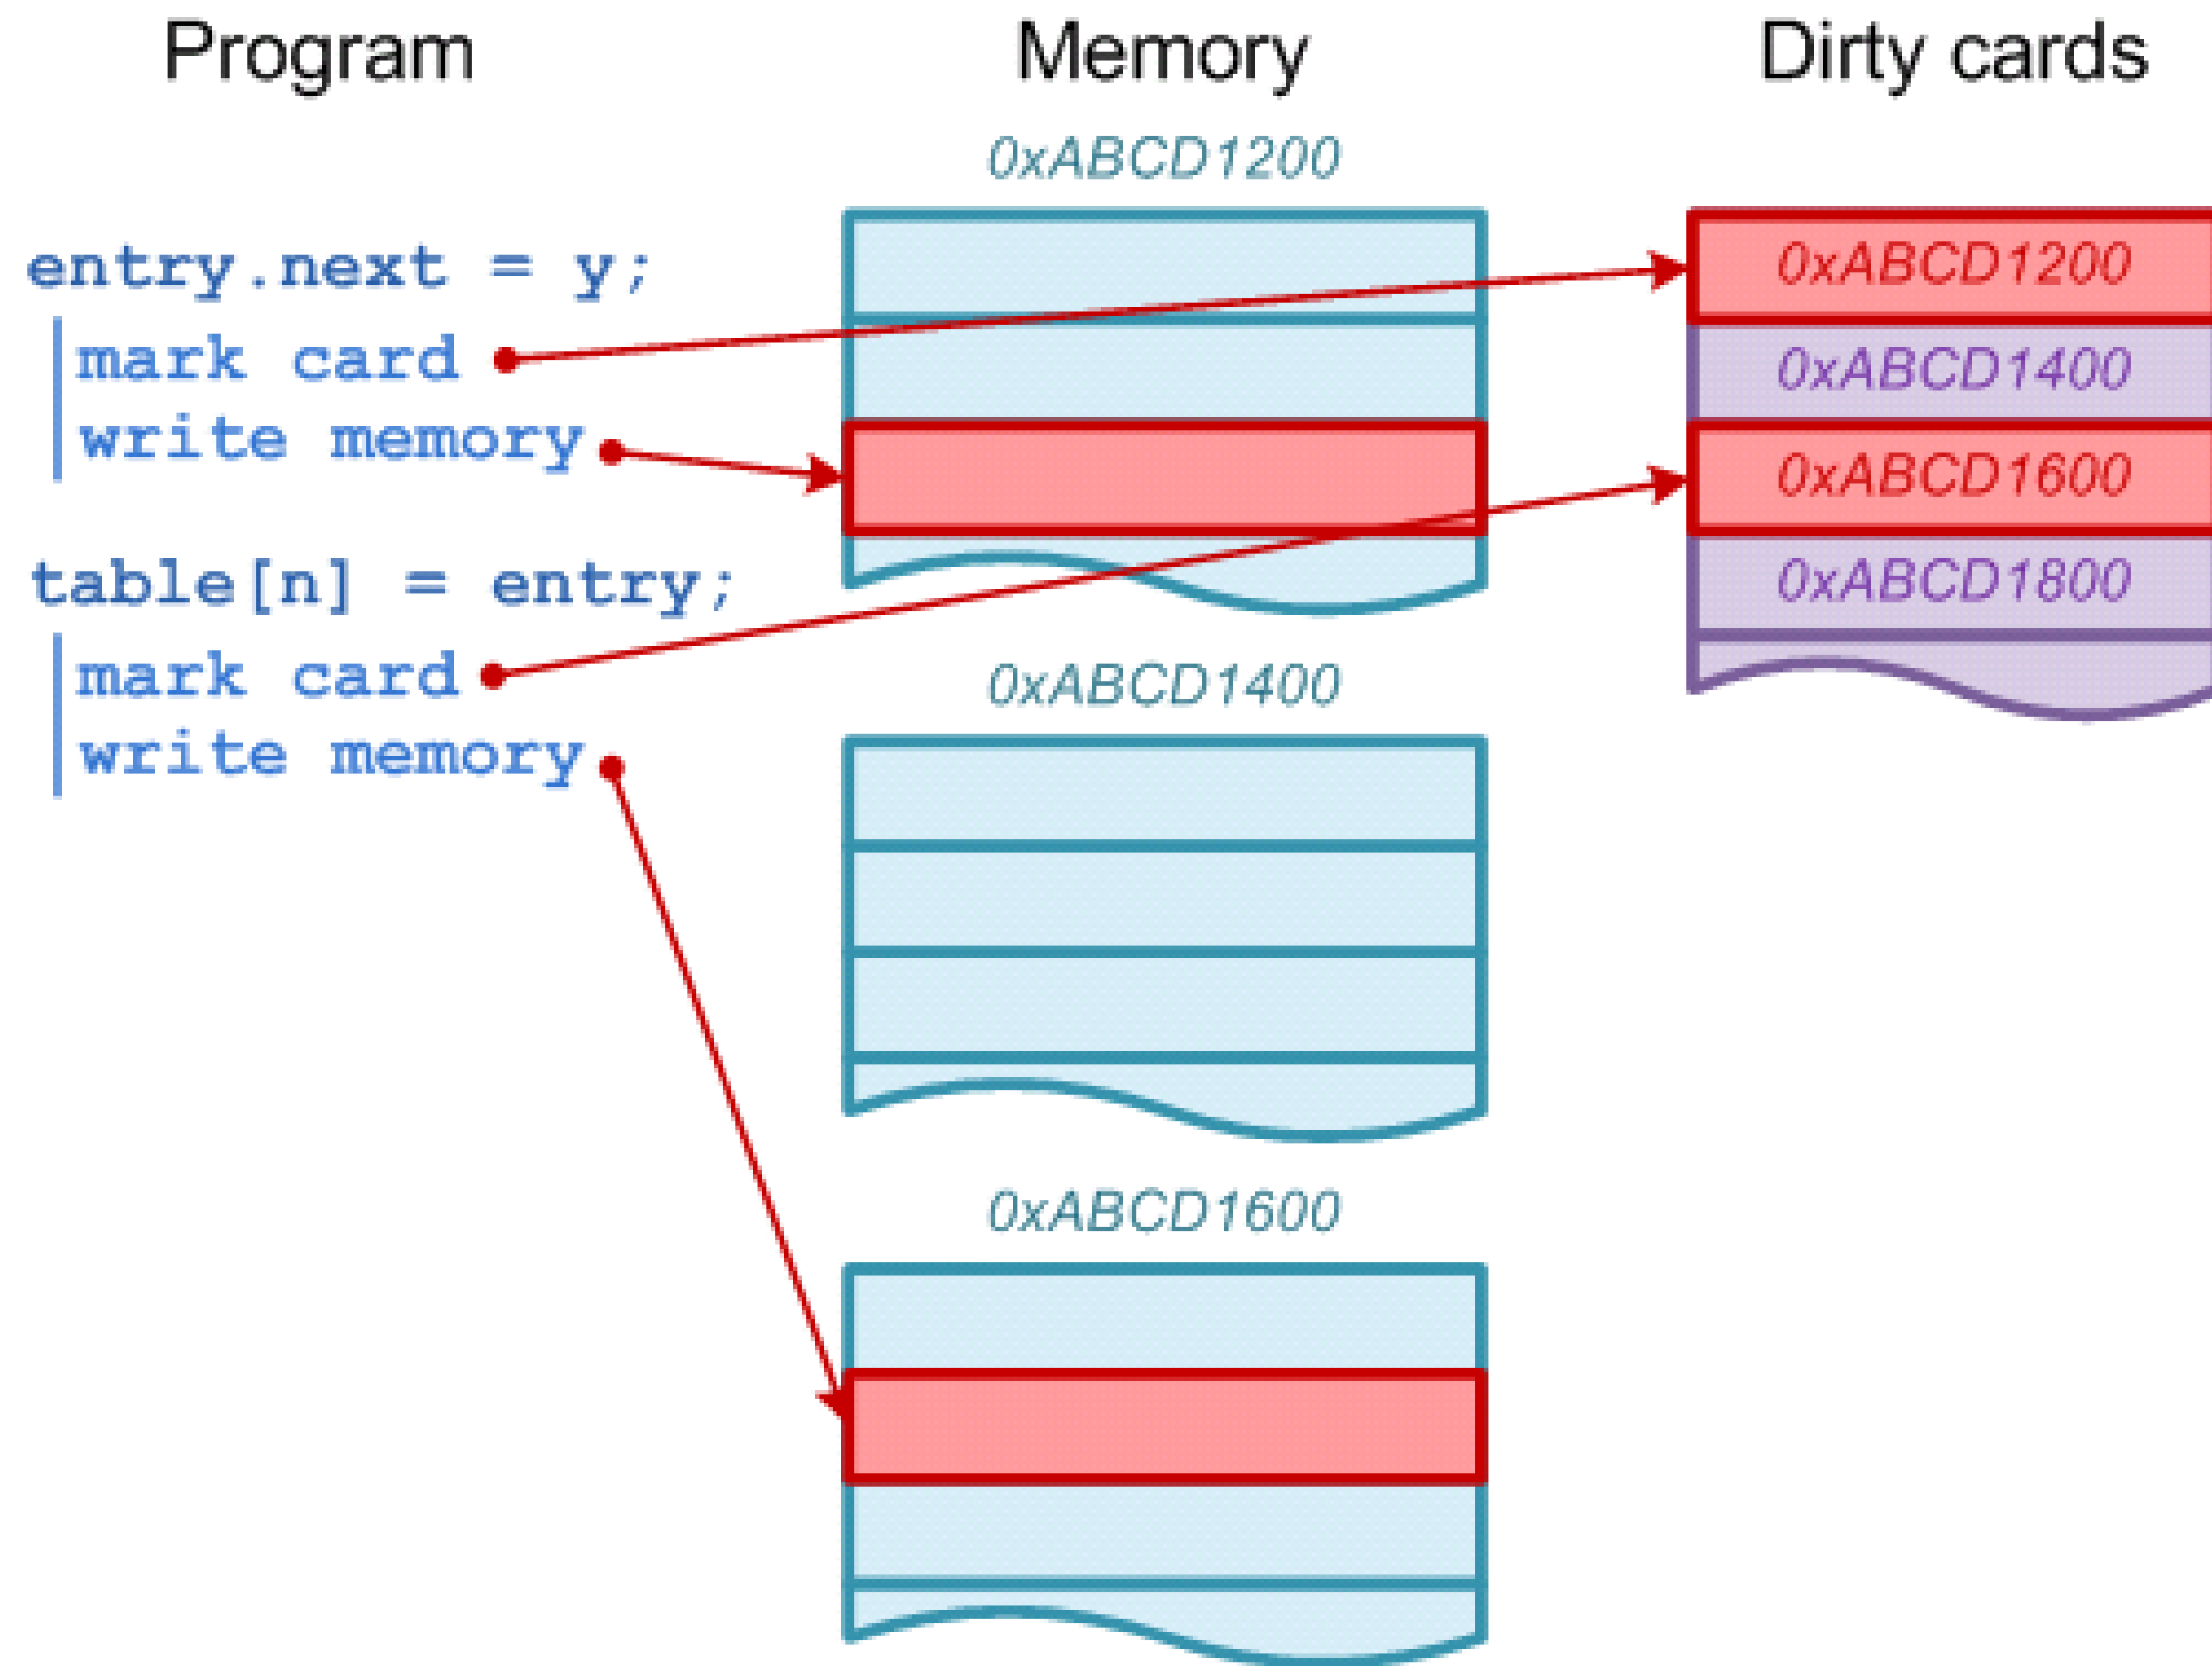
- Must be decompressed before read 'through'
- Can be copied without decompression
- Can be compared without decompression
- Balance size vs. computation

# Card Marking

“The JVM maintains a card map, with one byte corresponding to each card in the heap. Each time a pointer field in an object in the heap is modified, the corresponding byte in the card map for that card is set.”



Brian Goetz - <http://www.ibm.com/developerworks/library/j-jtp11253/>



# CardMarking instanceof WriteBarrier

- An optimisation for young collections
- Reduce the impact of OldGen size on scan time
- Introduce a small overhead
- Comes in different flavours!



# CardMarking v1 (default)

*; rsi is 'this' address*

*; rdx is setter param, reference to bar*

*; this.foo = bar*

**mov QWORD PTR [rsi+0x20],rdx ; <- WHAT WE WANTED**

**mov r10,rsi**

*; r10 = r10 >> 9;*

**shr r10,0x9**

*; r11 is base of card table, byte[] CARD\_TABLE*

**mov r11,0x7ebdfcff7f00**

*; Mark 'this' card as dirty*

*; CARD\_TABLE[this address >> 9] = 0*

**mov BYTE PTR [r11+r10\*1],0x0**

# CardMarking v2

## (-XX:+UseCondCardMark)

*; rsi is 'this' address*

*; rdx is setter param, reference to bar*

*; r10 = this*

**mov r10,rsi** <- WHAT WE WANTED

*; r10 = r10 >> 9*

**shr r10,0x9**

*; r11 = CARD\_TABLE*

**mov r11,0x7f7cb98f7000**

*; r11 = CARD\_TABLE + (this >> 9)*

**add r11,r10**

*; r8d = CARD\_TABLE[this >> 9]*

**movsx r8d,BYTE PTR [r11]**

**test r8d,r8d**

*; if(CARD\_TABLE[this >> 9] == 0) goto 0x00007fc4a1071d7d*

**je 0x00007fc4a1071d7d**

*; CARD\_TABLE[this >> 9] = 0*

**mov BYTE PTR [r11],0x0**

0x00007fc4a1071d7d:

**mov QWORD PTR [rsi+0x20],rdx** ; this.foo = bar

# CardMarking v3

## (-XX:+UseG1GC)

```
movsx edi,BYTE PTR [r15+0x2d0] ; read GC flag
cmp edi,0x0; if (flag != 0)
jne 0x00000001066fc601; GOTO OldValBarrier
Label WRITE:
```

```
mov QWORD PTR [rsi+0x20],rdx; this.foo = bar <- WHAT WE WANTED
mov rdi,rsi; rdi = this
xor rdi,rdx; rdi = this XOR bar
shr rdi,0x14; rdi = (this XOR bar) >> 20
cmp rdi,0x0; If this and bar are not same gen
jne 0x00000001066fc616; GOTO NewValBarrier
Label EXIT:
```

```
;...
Label OldValBarrier:
mov rdi,QWORD PTR [rsi+0x20]
cmp rdi,0x0; if(this.foo == null)
je 0x00000001066fc5dd; GOTO WRITE
mov QWORD PTR [rsp],rdi ; setup rdi as parameter
call 0x000000010664bca0 ; {runtime_call}
jmp 0x00000001066fc5dd; GOTO WRITE
```

```
Label NewValBarrier:
cmp rdx,0x0; bar == null
je 0x00000001066fc5f5 goto Exit
mov QWORD PTR [rsp],rsi
call 0x000000010664bda0 ; {runtime_call}
jmp 0x00000001066fc5f5 ; GOTO exit;
```

# So What?

- References mean extra work
- Impact can change by option/GC/JVM
- ‘Normalized/Flat’ data structures can help
  - Inheritance vs. Composition?
  - Consider access patterns
- Value Types might help (Java 9?)
- ObjectLayout might help (Java 9?)



```
while (hasQ() && hasTime()){  
    tryA();  
}  
return;
```